

code for checking the Quality of Wine

```
In [3]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import plotly.express as px
import seaborn as sns
import matplotlib.pyplot as plt
import plotly.graph_objects as go
from plotly.subplots import make_subplots

from xgboost import XGBClassifier,XGBRegressor

In [12]: import os
for dirname, _, filenames in os.walk("C:\\\\Users\\amith\\Downloads\\winequality-red.csv"):
    for filename in filenames:
        print(os.path.join(dirname, filename))

In [10]: # Importing the necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

In [11]: # Load the dataset (winequality-red.csv or winequality-white.csv)
url = 'https://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-red.csv'
data = pd.read_csv(url, sep=';')

In [12]: # Display the first 5 rows of the dataset
print(data.head())

    fixed acidity  volatile acidity  citric acid  residual sugar  chlorides  \
0              7.4              0.70         0.00           1.9         0.076
1              7.8              0.88         0.00           2.6         0.098
2              7.8              0.76         0.04           2.3         0.092
3             11.2              0.28         0.56           1.9         0.075
4              7.4              0.70         0.00           1.9         0.076

    free sulfur dioxide  total sulfur dioxide  density  pH  sulphates  \
0                11.0                34.0  0.9978  3.51         0.56
1                25.0                67.0  0.9968  3.20         0.68
2                15.0                54.0  0.9970  3.26         0.65
3                17.0                60.0  0.9980  3.16         0.58
4                11.0                34.0  0.9978  3.51         0.56

    alcohol  quality
0         9.4         5
1         9.8         5
2         9.8         5
3         9.8         6
4         9.4         5

In [13]: # Data overview
print(data.info())
print(data.describe())

# Check for missing values
print(data.isnull().sum())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1599 non-null  float64
1   volatile acidity       1599 non-null  float64
2   citric acid            1599 non-null  float64
3   residual sugar         1599 non-null  float64
4   chlorides              1599 non-null  float64
5   free sulfur dioxide     1599 non-null  float64
6   total sulfur dioxide    1599 non-null  float64
7   density                1599 non-null  float64
8   pH                    1599 non-null  float64
9   sulphates              1599 non-null  float64
10  alcohol                1599 non-null  float64
11  quality                1599 non-null  int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
None

      fixed acidity  volatile acidity  citric acid  residual sugar  \
count    1599.000000      1599.000000    1599.000000      1599.000000
mean         8.319637         0.527821     0.270976         2.538806
std         1.741096         0.179060     0.194801         1.409928
min         4.600000         0.120000     0.000000         0.900000
25%         7.100000         0.390000     0.090000         1.900000
50%         7.900000         0.520000     0.260000         2.200000
75%         9.200000         0.640000     0.420000         2.600000
max        15.900000         1.580000     1.000000        15.500000

      chlorides  free sulfur dioxide  total sulfur dioxide  density  \
count    1599.000000      1599.000000      1599.000000    1599.000000
mean         0.087467         15.874922         46.467792     0.996747
std         0.047065         10.460157         32.895324     0.001887
min         0.012000         1.000000         6.000000     0.990070
25%         0.070000         7.000000        22.000000     0.995600
50%         0.079000        14.000000        38.000000     0.996750
75%         0.090000        21.000000        62.000000     0.997835
max         0.611000        72.000000       289.000000     1.003690

      pH  sulphates  alcohol  quality
count    1599.000000    1599.000000    1599.000000    1599.000000
mean         3.311113         0.658149     10.422983         5.636023
std         0.154386         0.169507         1.065668         0.807569
min         2.740000         0.330000         8.400000         3.000000
25%         3.210000         0.550000         9.500000         5.000000
50%         3.310000         0.620000        10.200000         6.000000
75%         3.400000         0.730000        11.100000         6.000000
max         4.010000         2.000000        14.900000         8.000000
fixed acidity          0
volatile acidity       0
citric acid            0
residual sugar         0
chlorides              0
free sulfur dioxide    0
total sulfur dioxide   0
density                0
pH                    0
sulphates              0
alcohol                0
quality                0
dtype: int64

In [14]: # Split data into features (X) and target (y)
X = data.drop('quality', axis=1)
y = data['quality']

In [18]: # Split the dataset into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

In [19]: from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import cross_val_score

# Initialize Random Forest model
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)

# Train the model
rf_model.fit(X_train, y_train)

# Predict on test set
y_pred_rf = rf_model.predict(X_test)

In [20]: # Evaluate the model
rf_r2 = r2_score(y_test, y_pred_rf)
rf_mse = mean_squared_error(y_test, y_pred_rf)
rf_rmse = np.sqrt(rf_mse)

print(f'Random Forest R² Score: {rf_r2}')
print(f'Random Forest MSE: {rf_mse}')
print(f'Random Forest RMSE: {rf_rmse}')

Random Forest R² Score: 0.5390429623873638
Random Forest MSE: 0.3012381249999997
Random Forest RMSE: 0.5488516420673258

In [23]: # Feature importance analysis
coefficients = pd.DataFrame(model.coef_, X.columns, columns=['Coefficient'])
print(coefficients)

      Coefficient
fixed acidity    0.022593
volatile acidity -1.042144
citric acid      -0.103974
residual sugar   0.013310
chlorides        -1.764144
free sulfur dioxide  0.004627
total sulfur dioxide -0.003248
density          -17.239629
pH               -0.342743
sulphates        0.876382
alcohol          0.271549

In [24]: # Assign values for the features (as a single data point)
# Example: Create a dictionary with custom feature values
new_data = {
    'fixed acidity': 7.4,
    'volatile acidity': 0.7,
    'citric acid': 0.0,
    'residual sugar': 1.9,
    'chlorides': 0.076,
    'free sulfur dioxide': 11.0,
    'total sulfur dioxide': 34.0,
    'density': 0.9978,
    'pH': 3.51,
    'sulphates': 0.56,
    'alcohol': 9.4
}

# Convert the dictionary to a DataFrame (single row of feature values)
new_data_df = pd.DataFrame([new_data])

# Display the new data point
print("New data point:")
print(new_data_df)

# Use the trained model to predict wine quality for the new data point
predicted_quality = model.predict(new_data_df)

# Display the predicted quality
print(f'\nPredicted Wine Quality: {predicted_quality[0]}')

New data point:
      fixed acidity  volatile acidity  citric acid  residual sugar  chlorides  \
0              7.4              0.7           0.0           1.9         0.076

      free sulfur dioxide  total sulfur dioxide  density  pH  sulphates  \
0                11.0                34.0  0.9978  3.51         0.56
```

alcohol  
0 9.4

Predicted Wine Quality: 5.040537396744657